

Intro to Great Lakes Tutorial

This tutorial will lead you through a basic introduction of transferring data to Great Lakes and running a simple machine learning project as a batch submission job. All code and data will be provided.

Prerequisites

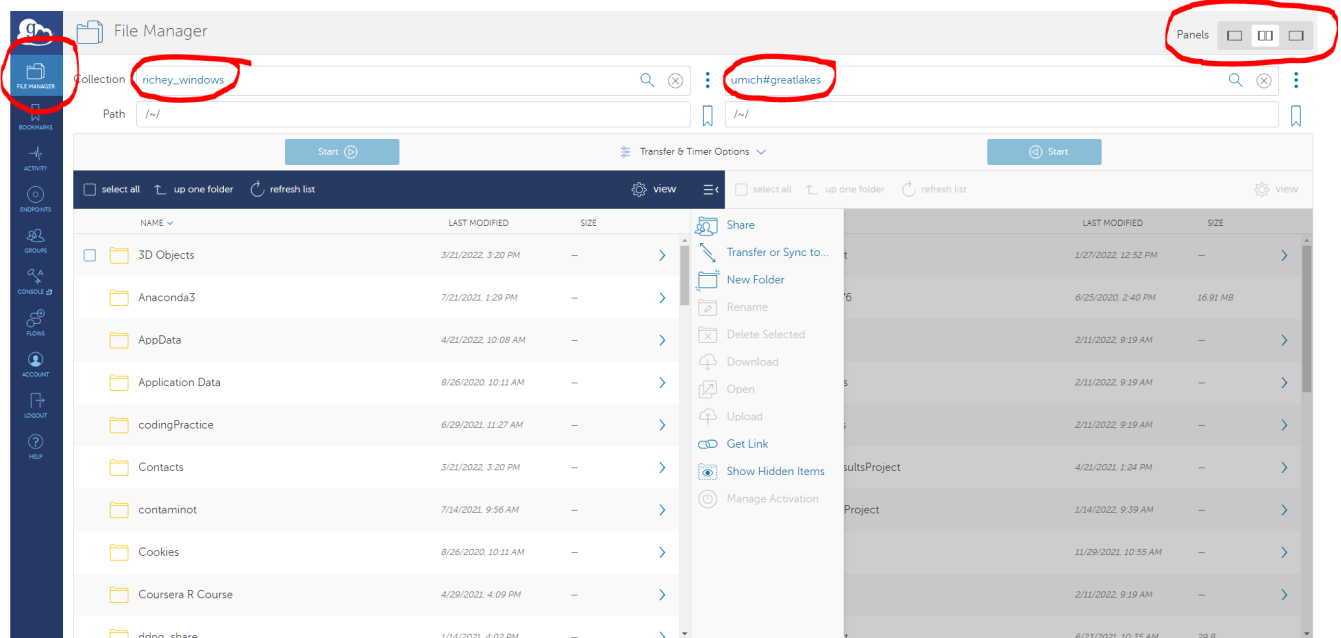
- Access to Great Lakes
- Familiarity with Python
- Familiarity with Github

Tutorial Outcomes

- Be able to transfer data from personal computer to Great Lakes using Globus
- Know how to store and access data in /scratch
- Be familiar with job scripts and how to set different parameters
- Load modules for a project
- Submit a batch job and analyze outputs

Process

1. Download the github repository containing the code for this tutorial at <https://github.com/richey-m-umich/greatlakes-tutorials>
2. Install and set up a Globus Endpoint following [these instructions](#) for your appropriate operating system.
3. Log into globus.org using your umich credentials. Set up your opening page according to the following:
 - a. Choose "File Manager" in the left toolbar
 - b. Choose "two panels" in the top right corner
 - c. On the left panel, navigate to your personal computer endpoint in the collection name
 - d. On the right panel, enter "umich#greatlakes" in the collection name
 - e. Your screen will now look like the following screenshot



4. Log onto Great Lakes using the command line
5. Create a new folder inside **/scratch/<account>** called **tutorials/intro-to-greatlakes**
6. Create a new folder in account home directory called **tutorials**
7. Transfer data to Great Lakes
 - a. In the left panel on Globus, navigate to the location of the downloaded code
 - b. In the right panel, in the Path line, type **/scratch/<account>/tutorials/intro-to-greatlakes** and navigate to that location
 - c. In the left panel, navigate into the **data** folder and select **iris.csv**
 - d. On the top of the left panel, press **'Start'**
8. Transfer code to Great Lakes
 - a. In the left panel on Globus, navigate to the location of the downloaded code
 - b. In the right panel, return to the home folder and enter the workshops folder by typing **/~/tutorials/**
 - c. In the left panel, select **intro-to-greatlakes** (entire folder)
 - d. On the top of the left panel, press **'Start'**
9. Modify training and validation scripts
 - a. Navigate to **intro-to-greatlakes** in your command line and open **training_script.py**
 - b. On line 13, change the location of the input file to your local scratch directory
 - c. Navigate to **intro-to-greatlakes** in your command line and open **validation_script.py**
 - d. On line 9, change the location of the input file to your local scratch directory
10. Modify job training script
 - a. Navigate to **intro-to-greatlakes** and open **submit_training.sh**
 - b. Change the account to your account name (if staff, can use hpcstaff)
 - c. Change job name to **submit-training-<username>**
 - d. Change email to own email

- e. Save file
11. Modify analysis job script
 - a. Navigate to **intro-to-greatlakes** and open **submit_validation.sh**
 - b. Change account to your account name (if staff, can use hpcstaff)
 - c. Change job name to **submit-validation-<username>**
 - d. Change email to own email
 - e. Save file
12. Submit training job script on compute node
 - a. Navigate to ml_on_greatlakes
 - b. Type **sbatch submit_training.sh**
 - c. Output will be written to **slurm-<job number>.out**
 - d. To check the status of your running jobs, type **squeue -u <username>**
13. Submit validation job script on compute node
 - a. Navigate to ml_on_greatlakes
 - b. Type sbatch **submit_validation.sh**
 - c. Output will be written to **slurm-<job number>.out**

Notes and Options

Configuration Options

- **--account:** Where to charge the job to
 - This is different than your user login!
- **--job-name:** Easily identifiable string to refer to the job
- **--mail-user:** Where to send information about the job
- **--mail-type:** How often to send information about the job
 - Most common are BEGIN and END
 - Others are NONE, FAIL, TIME_LIMIT, ALL, etc.
- **--nodes:** How many nodes from the cluster to use to complete the job
 - If not specified, defaults to allocate enough for the job
 - Using more nodes can help speed up parallel processing
- **--ntasks-per-node:** How many tasks to get sufficient resources
 - Default is one task per node
- **--cpus-per-task:** How many cpus on each node to allocate
 - If not specified, defaults to one processor per task
- **--mem-per-cpu:** Minimum memory required per allocated CPU
 - Generally only want to change if working with large memory tasks
- **--time:** Limit on total run time of the job
 - Can use this to help manage costs
- **--partition:** Which queue to allocate the job to
 - Options are standard (default), gpu (GPU jobs), largemem (large memory jobs), viz, debug, standard-oc (on-campus software)

Helpful Commands

- **squeue -u <username>** will show status of queued jobs
- **scancel <job number>** will cancel a queued job

Conclusion

By the end of this tutorial, you are now able to run a batch job on Great Lakes and evaluate the output. Please contact Meghan Dailey (richeym@umich.edu) for any questions or updates to this tutorial.